


DS-UA 301
Advanced Topics in Data Science
*Advanced Techniques in ML and Deep
Learning*

LECTURE 5
Parijat Dube


Deep Learning for Computer Vision

Image classification


airplane




automobile



bird




cat




deer




dog




frog




horse



ship



truck




Other vision tasks

Classification




Cougar

Classification + Localization




Cougar

Object Detection




Cougar, Butterfly

Pixel and their values




The range of intensity values from 0 (black) to 255 (white).



150 154 160 157 106 140 147 142 141 147 132 150 171 117 136
144 159 125 121 157 143 132 136 153 138 155 164 169 162 152
190 175 169 155 161 136 152 158 141 162 147 153 161 168 169
185 203 139 161 151 159 145 167 179 167 150 155 165 159 158
151 153 163 152 160 152 164 131 131 51 124 152 154 145 143
164 162 158 167 157 164 166 139 132 138 119 148 154 139 146
147 148 143 155 169 160 152 161 159 143 138 163 132 152 146
66 129 163 165 163 161 154 157 167 162 174 153 156 151 156
162 173 172 161 158 158 159 167 171 169 164 159 158 159 162
163 164 161 155 155 158 161 167 171 168 162 162 163 164 166
167 167 165 163 160 160 164 166 169 168 164 163 165 167 170
172 171 170 170 166 163 166 170 169 168 167 165 163 163 160
173 172 170 169 166 163 167 169 170 170 170 165 160 157 148
167 168 165 173 173 172 167 170 170 171 171 169 162 163 162
200 198 189 196 191 188 163 168 172 177 177 186 180 180 188

<http://whydomath.org/node/wavlets/imagebasics.html>

RGB Colorspace



Need for special neural network architectures

- **Fully-Connected Neural Networks** for image recognition/classification:

1. **Dimension Problem:**


- Reshaping to 1D: losing the spatial structure

2. **Size Problem:**

- Example features: $1000 \times 1000 \times 3$ [Y \times X \times RGB]
- Weights: 3M for each neuron just for first layer
- 100 layers: about 300M DP weights => **2.4GB model size**


3. **Overfitting problem:**

- One weight per pixel leads to **overfitting**




Which Architecture?

- Convolutional neural networks have three main traits that support their use as models of biological vision: (1) they can perform visual tasks at near-human levels, (2) they do this with an architecture that replicates basic features known about the visual system, and (3) they produce activity that is directly relatable to the activity of different areas in the visual system.



<https://gracewlindsay.com/2018/05/17/deep-convolutional-neural-networks-as-models-of-the-visual-system-qa/>



Convolutional Neural Networks (CNNs)

- Widely used in computer vision
- Specialized structure: inspired by Hubel and Wiesel experiments on visual cortexes of cat and monkey
- Success in ImageNet (ILSVRC) competitions brought them into limelight since 2012
- Layers have *volume*: length, width, depth
- The depth is 3 for color images, 1 for grayscale, and an arbitrary value for hidden layers.
- Three basic operations in CNNs are convolution, pooling, and ReLU.
 - The convolution operation is analogous to the matrix multiplication in a conventional network


How CNNs work?

- <https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-from-scratch-for-mnist-handwritten-digit-classification/>

Example of MNIST classification

- https://python-course.eu/neural_network_mnist.php


Example of Convolution



- Add up over multiple activations maps from the depth

Convolution Layer with 3 dimensions

32x32x3 image




5x5x3 filter




- Preserve the spatial structure of the image
- Convolve the filter with the image i.e. “slide over the image spatially, computing dot products”
- Filters always extend the full depth of the input volume (the 3 channels in this example)

Output image: Activation Map




Multiple filters Activation Map

- If we had 6 5x5x3 filters, we'll get **6 separate activation maps**
- We stack these up to get a “new image” of size 28x28x6!





Convolutional Network

- A Convolution Network is a sequence of Convolutional Layers, interspersed with activation functions



Convolution Layer: Size of output

$L_q \times B_q \times d_q$  $F_q \times F_q \times d_q$ $(L_q - F_q + 1) \times (B_q - F_q + 1)$	<div style="text-align: left; margin-left: 20px;">with stride S_q</div> $L_q \times B_q \times d_q$  $F_q \times F_q \times d_q$ $(L_q - F_q) / S_q + 1 \times (B_q - F_q) / S_q + 1$
---	--

When a stride of S_q is used in the q th layer, the convolution is performed at the locations $1, S_q + 1, 2S_q + 1$, and so on along both spatial dimensions of the layer.

Convolution Neural Networks Properties

- **Sparse connectivity** because we are creating a feature from a region in the input volume of the size of the filter (usually much smaller than the input size)
 - Each neuron in the output layer has connectivity to a small subset of neurons in the input layer which are spatially close
- **Shared weights** because we use the same filter across entire spatial volume
 - Interpret a shape in various parts of the image in the same way
 - Reduces the number of parameters to learn
- A feature in a hidden layer captures some properties of a region of the input image.

Padding

- Convolution

- Image size 6×6 ($n = 6$)
- Filter size $f = 3 \times 3$ ($f = 3$)
- Output: 4×4 ($n' = 4$)
- Output size formula:
 - size: $n' \times n'$
 - where $n' = n - f + 1$

6 x 6 image

8	8	8	0	0	0
8	8	8	0	0	0
8	8	8	0	0	0
8	8	8	0	0	0
8	8	8	0	0	0
8	8	8	0	0	0

dot

3 x 3 filter

1	0	-1
1	0	-1
1	0	-1

=

4 x 4 image


0	24	24	0
0	24	24	0
0	24	24	0
0	24	24	0

- Problem:

- Edges are used less by the convolution (less relevant)
- Output image shrinks

Padding - Definition

- Padding:
 - Extend input image adding a “frame”
- Convolution size:
 - Input image: $n \times n$
 - Filter: $f \times f$
 - Output image: $n' \times n'$
 - **Without padding:**
 - $n' = n - f + 1$
 - **With padding of p :**
 - $n' = n + 2p - f + 1$
- More definitions:
 - “Same” convolution: convolution with padding when output image and input image size are the same
 - “Valid” convolution: no padding
- In PyTorch you have to specify the size of the padding (Default is 0)

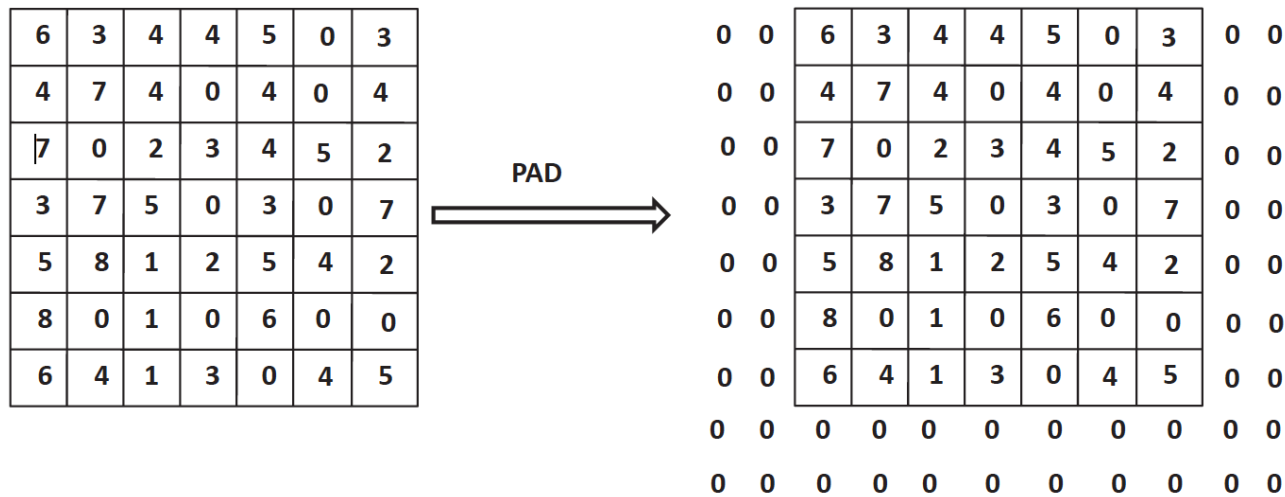


- Example Convolution with padding of 1:
 1. Do padding and obtain a 8 x 8 image
 2. Do convolution and obtain a 6 x 6 image

Padding

- Convolution results in loss of information
 - Size of the output layer is smaller than the size of the input layer
 - Loss of information along the borders of the image (or of the feature map, in the case of hidden layers).
- By adding $(F_q - 1)/2$ “pixels” all around the borders of the feature map, one can maintain the size of the spatial image by taking stride =1
- Padding can be of any size up to $F_q - 1$
- What happens if you add a padding of size F_q or greater ?

Padding example



- A padding of size P increases both the length and width of input by 2P


Stride

- Sometimes we want to obtain a **smaller output image** than the standard convolution or we want to have a lighter computation
- We can have the filter **jump a few pixels** at each step (stride)
- There are several implications:
 - The **compute time will diminish**
 - The **output size also will diminish**
 - The **convolution will be less fine-grained** (precise)

Stride - Definition

- Stride:
 - Jump S elements when moving filter
 - If filter falls outside image:
 - ignore computation: no output
- Convolution size:
 - Input image: $n \times n$
 - Filter: $f \times f$
 - Output image: $n' \times n'$
 - Stride size: s
 - **With padding** of p and stride s :

$$n' = \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$




- Example Convolution with stride of 2
 - Input image: 7x7
 - Filter: 3x3
 - Stride: 2
 - Output image $n' \times n'$: 3x3

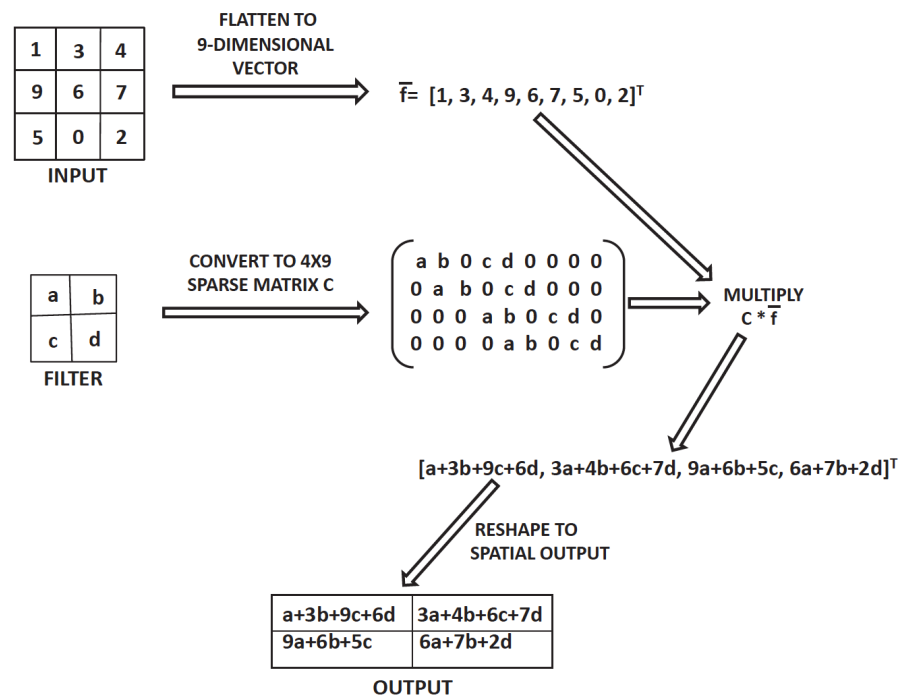
$$n' = \left\lfloor \frac{7 + 0 - 3}{2} + 1 \right\rfloor = 3$$

Max Pooling

- Pooling operates on square regions of size $P \times P$ in each of the activation maps of the input and returns maximum of the values in those regions
- Pooling preserves the depth; depth of output is same as input
- Pooling is a subsampling technique; Greatly reduces the size of output activation maps and hence the number of parameters to learn in subsequent layers



Convolution as a Matrix Operation




Convolution Summary

Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.


Convolution example – Image Classification

- Says whether a picture is of an X or an O





Feature recognition Example

1. Line up the feature and the image patch.
2. Multiply each image pixel by the corresponding feature pixel.
3. Add them up.
4. Divide by the total number of pixels in the feature [**optional: to normalize values**]



Feature recognition



Feature recognition


1	-1	-1
-1	1	-1
-1	-1	1

1	1	1
1	1	1
1	1	1

[illegible]

$$\frac{1+1+1+1+1+1+1+1+1}{9} = 1$$

Feature recognition



Convolution: Search every possible match

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1




1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

Convolution Layer


- Using Multiple Filters: One image becomes a stack of features (filtered images)



Pooling: Shrinking the image stack

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually the same window size, so there are no overlaps between windows).
3. Walk your window across your filtered images.
4. From each window, take the maximum (or average) value.

Max Pooling



Max Pooling

A stack of images becomes a stack of smaller images.

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33




0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

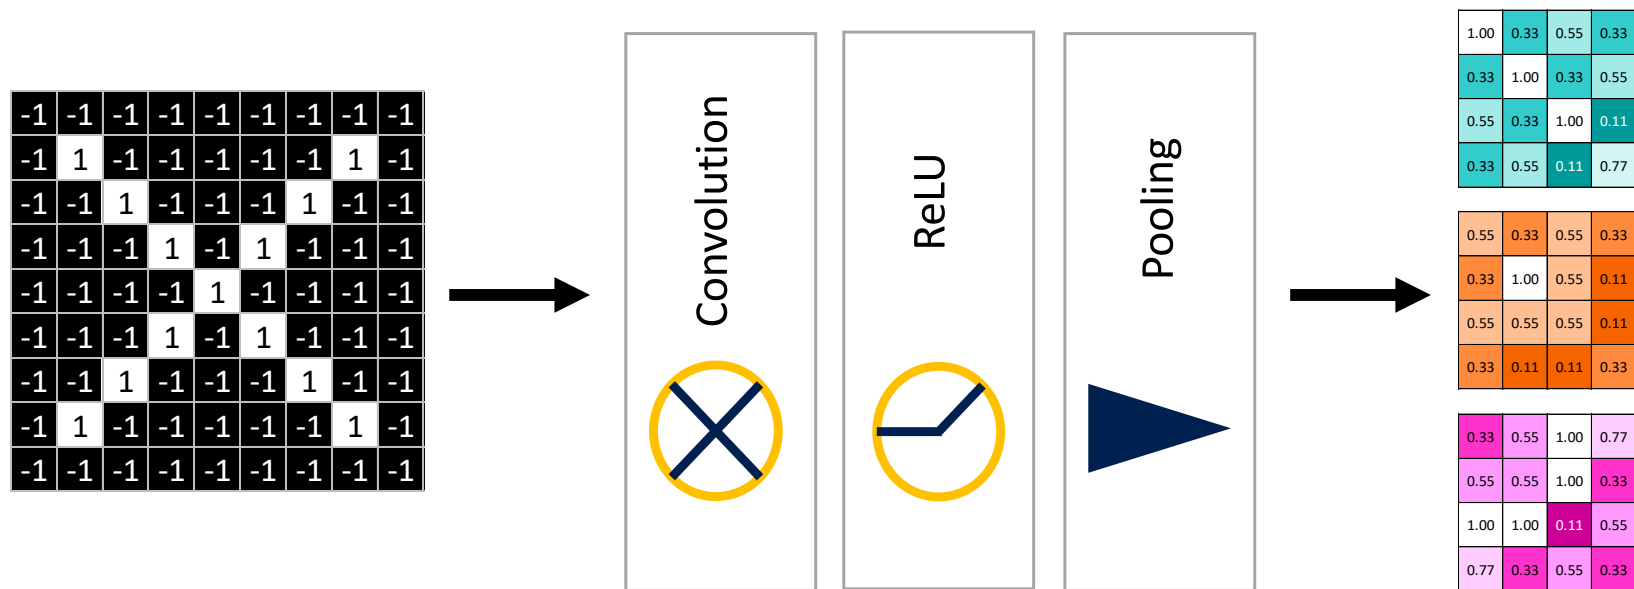


0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33


Rectified Linear Units (ReLUs)



Layers get stacked

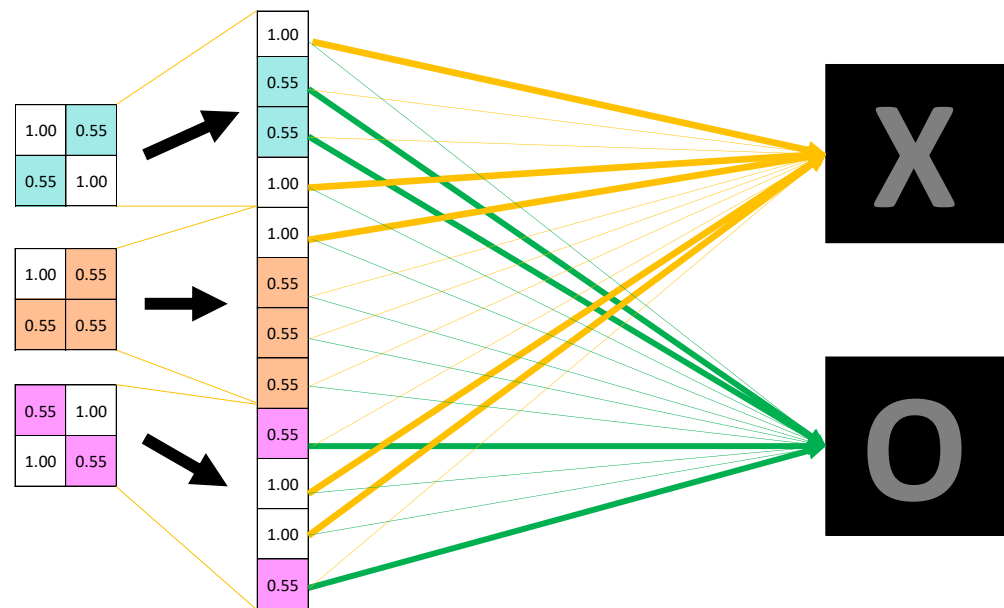


Layers get stacked



Classification with a Fully Connected Layer


- Stack of images obtained with multiple filters can be the input of FC layer



Creating Convolutional Neural Networks

- Interleaving of convolution, pooling, and ReLU layers
- ReLU often follows the convolutional layers
- Max-pooling is often applied after convolution-ReLU combinations
- LeNet-5 (convolutional network for hand-written digits classification)
 - 7 layers, 3 convolutional layers (C1, C3, C5), 2 sub-sampling (pooling) layers (S2 and S4), 1 fully connected layer (F6), and an output (O) layer
 - Did not use ReLU activation functions
 - Used (scaled) hyperbolic tangent as activation function in hidden layers
 - C5 is basically a fully connected layer as each unit is mapped to all the 16 input activation maps and filter size is same as input maps
 - Pattern: Conv→Pool→Conv→ Pool→FullyConn→ FullConn→ Output


LeNet-5



CNN for MNIST classification

- <https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-from-scratch-for-mnist-handwritten-digit-classification/>

CNN Application Example



- [Classifications of multispectral colorectal cancer tissues using convolution neural network, J Pathol Inform 2017, 8:1](#)